

UNIVERSITY OF HAMBURG

MASTER'S THESIS

A Global Analysis of Mesoscale Eddy Dynamics via a Surface-Signature-Based Tracking Algorithm

Author:

NIKOLAUS KOOPMANN

Co-Supervisor:

PROF. DR. JOHANNA
BAEHR

Supervisor:

PROF. DR. CARSTEN EDEN

2015/04/27

Abstract

Several global mesoscale ocean-eddy-dynamics censuses were obtained via an automated, sea-surface-anomaly-based detection- and tracking-algorithm. 12 years of input sea-surface-height data were taken from a satellite-altimetry product and from a global eddy-resolving ocean model. Variables of particular interest were horizontal eddy scale and zonal eddy drift speeds. Motivation was to answer the question whether time- and space-resolutions of the altimeter-product are sufficiently fine to resolve eddy scales accurately and to allow a successful tracking of individual eddies over time. The results suggest that the 0.25° -resolution of the merged satellite scans is insufficient to resolve realistic eddy scales in high latitudes and that a 7 day time-step to determine eddy drift-speeds is likely insufficient in regions of strong drift-vector gradients.

Contents

1 The Algorithm	5
1.1 Data Preparation	6
1.2 Rossby Radii and Phase Speeds	6
1.3 Find Contours	7
1.4 Filter Contours	8
1.5 Tracking	17
1.6 Running the Code	19
Appendix A Ideas for the Future	21
Appendix B 2-Dimensional Turbulence	25
Appendix C Eddy Categories	29
Appendix D Derivations	31
Appendix E Legend	35
Bibliography	39

1

The Algorithm

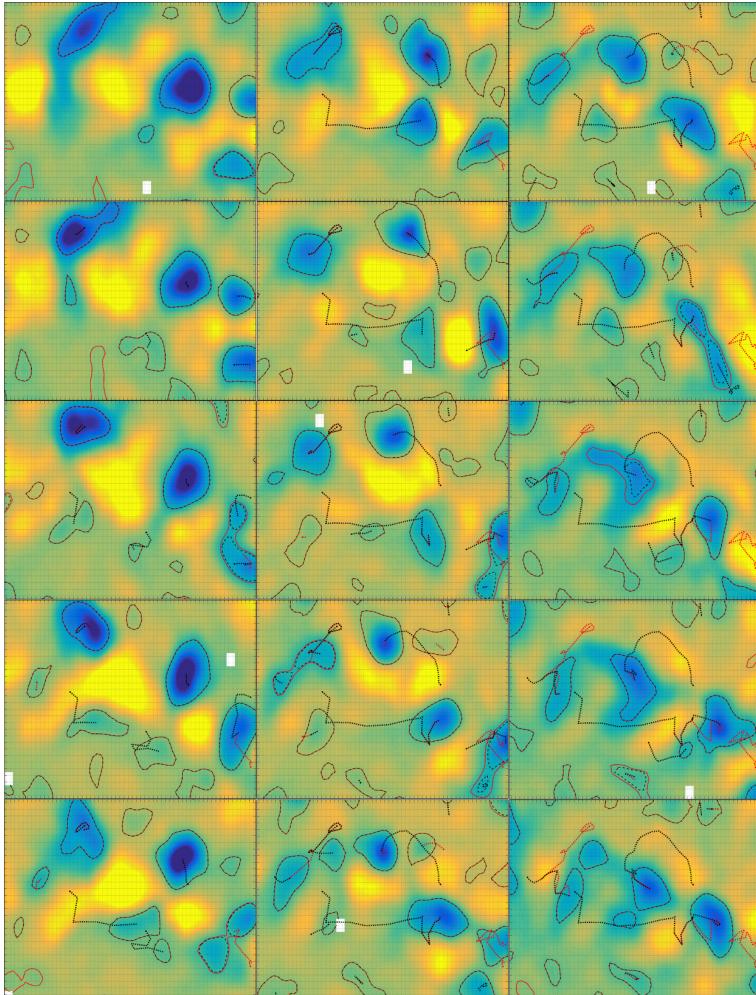


Figure 1.1: Fortnightly [Aviso-MII](#) snapshots (cyclones only) showing the area from 47.6°S to 30.1°S and 40.1°E to 64.9°E . Threshold age for saving is $8 \times 7\text{days}$. Red [black] color represents MI [MII]. Dashed lines are the contours and dotted lines are the tracks. Only *active* tracks are drawn. The general impression from animations of this sort is that the MI-method is good at tracking coherent, west-ward propagating, less-circular SSH anomalies while the MII-method seems superior at successfully tracking higher-amplitude vortices that get advected by mean currents (*e.g.* the strong cyclone in appr. the middle of the picture describing an anti-clockwise circular track due to advection by the ACC.)

THIS section walks through the algorithm step by step, so as to explain which methods are used and how they are implemented. The idea is that the code from step `s00..` on can only accept one particular structure of data. In earlier versions the approach was to write code that would adapt to different types of data automatically. All of this extra adaptivity turned out to

visually and structurally clog the code more than it did offer much of a benefit. The concept was therefore reversed. Input SSH-data needs to be altered to required format. Yet, there should be no need to adapt any of the later steps in any way. All input parameters are to be set in `INPUT.m` and `INPUTx.m`. A well-documented *core*-version (stripped of all experimental and redundant add-ons) will be made public on github presumably as early as end of June 2015 and no later than August 2015.

1.1 Step S00: Prepare Data

```
function S00_prep_data
```

BEFORE the actual eddy detection and tracking is performed, SSH-, latitude- and longitude-data is extracted from the given data at desired geo-coordinate bounds and saved as structures in the form needed by the next step (S01). This step also builds the file `window.mat` via `GetWindow3` which saves geometric information about the input and output data as well as a cross-referencing index-matrix which is used to reshape all *cuts* to the user-defined geo-coordinate-geometry. The code can handle geo-coordinate input that crosses the longitudinal seam of the input data. E.g. say the input data came in matrices that start and end on some (not necessarily strictly meridional) line straight across the Pacific and it is the Pacific only that is to be analyzed for eddies, the output maps are stitched accordingly. In the zonally continuous case *i.e.* the full-longitude case, an *overlap* in x-direction across the *seam*-meridian of the chosen map is included so that contours across the seam can be detected and tracked across it. One effect is that eddies in proximity to the seam can get detected twice at both zonal ends of the maps. The surplus double-eddies get filtered out in `S05_track_eddies`.

1.2 Step S01b: Find Mean Rossby Radii and Phase Speeds

```
function S01b_BruntVaisRossby
```

THIS function...

- – ...calculates the pressure $P(z, \phi)$ in order to...
- ...calculate the Brunt-Väisälä-Frequency according to

$$N^2(S, T, P, \phi) = -\frac{g(\phi)}{P} \frac{\partial \rho(S, T, P)}{\partial z}$$
 in order to...
- – ...integrate the Rossby-Radius $L_R^1 = \frac{1}{\pi f} \int_H N dz$ and ...
- apply the long-Rossby-Wave dispersion relation to
 found L_R^1 to estimate Rossby-Wave phase-speeds $c = -\frac{\beta}{k^2 + (1/Lr)^2} \approx -\beta L_R^{1/2}$

The 3-dimensional matrices (S and T) are cut zonally into slices which then get distributed to the threads. This allows for direct matrix operations for all calculations which would otherwise cause memory problems due to the immense sizes of the 3d-data ¹.

¹ E.g. the POP data has dimensions $42 \times 3600 \times 1800$.

Step S02: Calculate Geostrophic Parameters

```
function S02_infer_fields
```

THIS step reads the cut **SSH** data from `s00_prep_data` to perform 2 steps:

1. Calculate a mean over time of $SSH(y, x)$.
2. • use one of the files' geo-information to determine f , β and g .
 • calculate geostrophic fields from **SSH** gradients.
 • calculate deformation fields (vorticity, divergence, stretch and shear) via the fields supplied by the last step.
 • calculate O_w .
 • Subtract the mean from step 1 from each $SSH(t)$ to filter out persistent SSH-gradients e.g. across the Gulf-Stream.

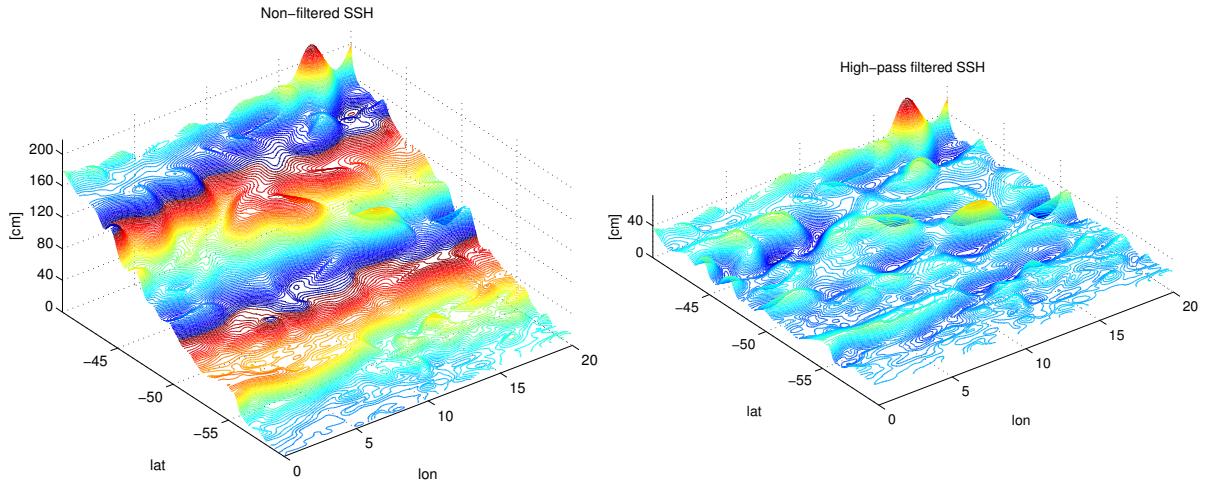


Figure 1.2: SSH with mean over time subtracted.

1.3 Step S03: Find Contours

```
function S03_contours
```

THE sole purpose of this step is to apply MATLAB's `contourc.m` function to the **SSH** data. It simply saves one file per time-step with all contour indices appended into one vector ². The

² see the MATLAB documentation.

contour intervals are determined by the user defined increment and range from the minimum- to the maximum of given **SSH** data.

The function `initialise.m`, which is called at the very beginning of every step, here has the purpose of rechecking the *cuts* for consistency and correcting the time-steps accordingly (*i.e.* when files are missing). `initialise.m` also distributes the files to the threads *i.e.* parallelization is in time dimension.

1.4 Step S04: Filter Contours

```
function S04_filter_eddies
```

SINCE indices of all possible contour lines at chosen levels are available at this point, it is now time to subject each and every contour to a myriad of tests to decide whether it qualifies for the outline of an eddy as defined by the user input threshold parameters.

1.4.1 Reshape for Filtering and Correct out of Bounds Values

```
function eddies2struct
function CleanEddies
```

In the first step the potential eddies are transformed to a more sensible format, that is, a structure `Eddies` of size `EddyCount`. The struct has fields for level, number of vertices, exact *i.e.* interpolated coordinates and rounded integer coordinates.

The interpolation of `contourc.m` sometimes creates indices that are either smaller than 0.5 or larger than ³ $N + 0.5$ for contours that lie along a boundary. After rounding, this seldomly leads to indices of either 0 or $N + 1$. These values get set to 1 and N respectively in this step.

³ where N is the domain size.

1.4.2 Descend/Ascend Water Column and Apply Tests

```
function walkThroughContsVertically
```

The concept of this step is a direct adaption of the algorithm described by Chelton *et al.* (2011). It is split into two steps, one for anti-cyclones and one for cyclones. Consider *e.g.* the anti-cyclone situation. Since all geostrophic anti-cyclones are regions of relative high pressure, all **ACs** effect an elevated **SSH** *i.e.* a *hill*. The algorithm ascends the full range of **SSH** levels where contours were found. Consider an approximately Gaussian shaped **AC** that has a peak **SSH** of say 5 increments larger than the average surrounding waters. As the algorithm approaches the sea surface from below, it will eventually run into contours that are closed onto themselves and that encompass the **AC**. At first these contours might be very large and

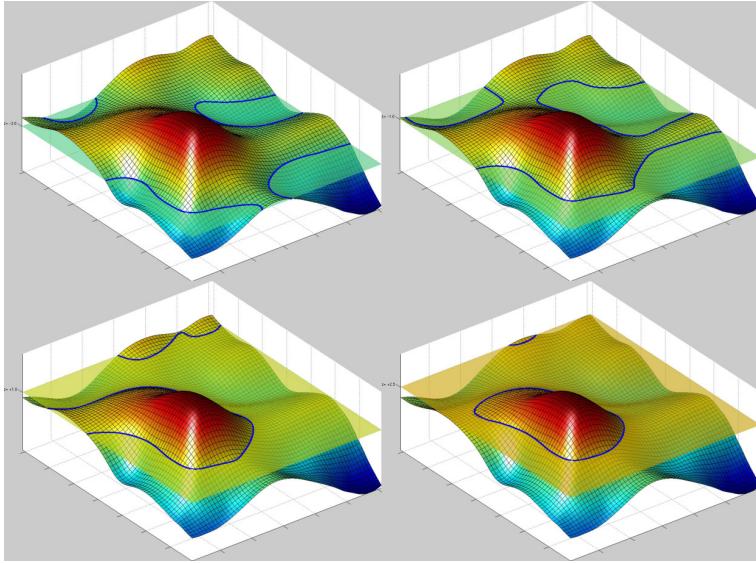


Figure 1.3: The algorithm approaches the appropriate vertical level incrementally.

describe not only one but several ACs and likely also cyclones, but as the algorithm continues upwards found contour will get increasingly circular, describing some outer *edge* of the AC . Once the contour and its interior pass all of the tests, the algorithm will decide that an AC was found and write it and all its parameters to disk. The AC 's region *i.e.* the interior of the contour will be flagged from here on. Hence any inner contour further up the water column will not pass the tests. Once all ACs are found for a given time-step, the SSH flags get reset and the entire procedure is repeated, only this time *descending* the SSH-range to find cyclones. The tests for cyclones and anti-cyclones are therefor identical except for a factor -1 where applicable. In the following the most important steps of the analysis are outlined.

Contour filter 1 NaN-Check Contour

```
function CR_RimNan
```

The first and most efficient test is to check whether indices of the contour are already flagged. Contours within an already found eddy get thereby rejected immediately.

Contour filter 2 Closed Ring

```
function CR_ClosedRing
```

Contours that do not close onto themselves are obviously not eligible for further testing.

Contour filter 3 Sub-Window

```
function get_window_limits, EddyCut_init
```

For further analysis a sub-domain around the eddy is cut out of the SSH data. These functions determine the indices of that window and subtract the resultant offset for the contour

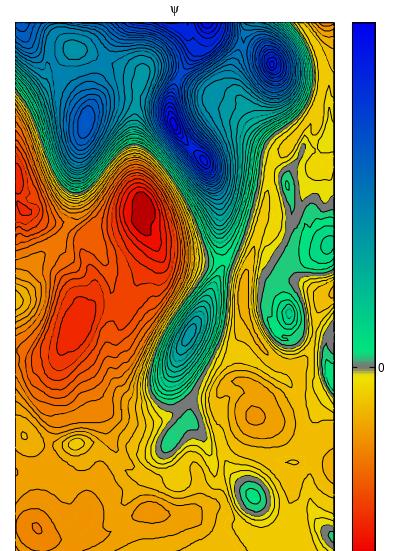


Figure 1.4: Stream function of a meandering jet shedding off a vortex. The line of strongest gradient *i.e.* fastest geostrophic speed later becomes the zero-vorticity-line at a theoretical distance σ from the center (Offset of Ψ is chosen arbitrarily).

indices.

Contour filter 4 Logical Mask of Eddy Interior

```
function EddyCut_mask
```

Basically this function creates a **flood-fill** logical mask of the eddy-interior. This is by far the most calculation-intensive part of the whole filtering procedure. A lot more time was wasted on attempting to solve this problem more efficiently than time could have been saved would said attempts have been successful. The current solution is basically just MATLAB's `imfill.m`, which was also used in the very first version of 09/2013. EDIT: `imfill.m` was replaced by using `inpoly.m` to determine which indices lie within the contour-polygon. This method seems to be more exact at determining whether the inside-part of one grid cell (with respect to the smooth, spline-interpolated contour) is larger than the outside part or not.

Contour filter 5 Sense

```
function CR_sense
```

All of the interior **SSH** values must lie either above or below current contour level, depending on whether anti-cyclones or cyclones are sought.

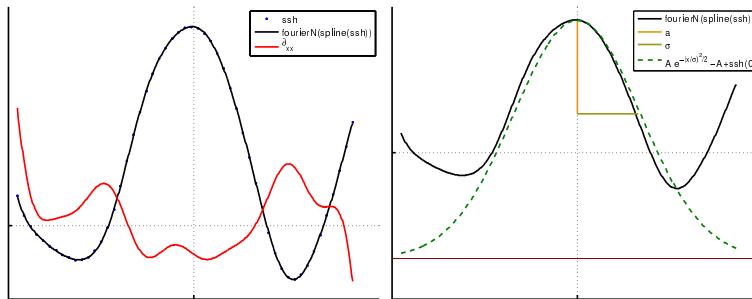


Figure 1.5: Left: Fourier-fit of an eddy from POP SSH-data and the 2nd derivative thereof. Right: Theoretical Gauss shape built from the resulting *standard-deviation i.e. σ* and amplitude.

Contour filter 6 Area

```
function getArea
```

The main goal here is to determine the area encompassed by the interpolated coordinates of the contour. It does so via MATLAB's `polyarea` function. This area is not related to the scale σ that is determined in contour-filter 12. It is however the relevant scale for the determination of the isoperimetric quotient in contour-filter 8.

If the respective switch is turned on, this function also checks that the area of found contour does not surpass a given threshold which in turn is a function of L_R^1 . Since L_R^1 gets very small in high latitudes a lower bound on the L_R^1 used here should be set as well. This is especially important for the southern ocean where L_R^1 gets very small while the strong mesoscale turbulence of the Antarctic circumpolar current

results in an abundance of relatively large eddies as far south as 60°S and beyond.

Contour filter 7 Circumference

`function EddyCircumference`

Circumference e.g. line-length described by the contour. This is the other parameter needed for contour-filter 8. This is however neither related to the actual eddy scale determined in contour-filter 12.

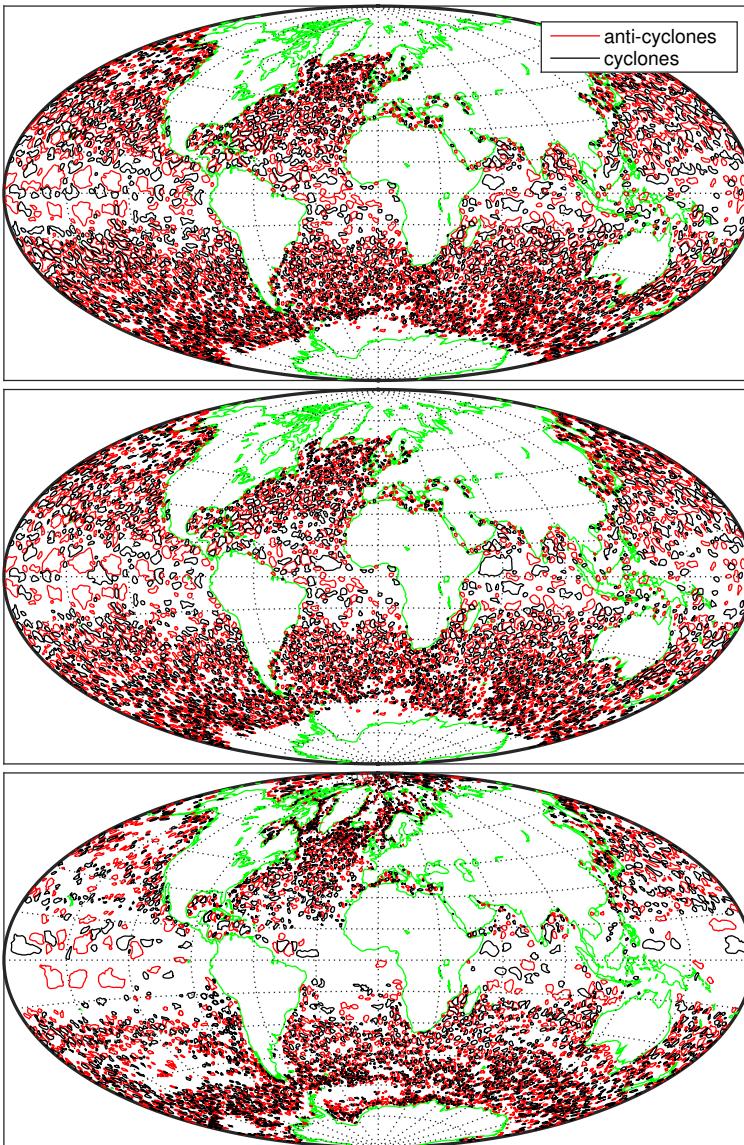


Figure 1.6: All contours that passed the filtering procedure for one exemplary time-step. Top: Aviso-MI . Mid: Aviso-MII . Bottom: POP-7day-MII .

Contour filter 8 Shape

`function CR_Shape`

This is the crucial part of deciding whether the object is *round enough*. A perfect vortex with $\frac{\partial u}{\partial y} = -\frac{\partial v}{\partial x}$ is necessarily a circle. The problem is that eddies get formed, die, merge, run into

obstacles, get asymmetrically advected etc. To successfully track them it is therefore necessary to allow less circle-like shapes whilst still avoiding to e.g. count 2 semi-merged eddies as one.

This is achieved by calculating the **isoperimetric quotient**, defined as the ratio of a ring's area to the area of a circle with equal circumference. Chelton *et al.* (2011) use a similar method. They require:

*The distance between any pair of points within the connected region must be less than a specified maximum (Chelton *et al.*, 2011).*

While this method clearly avoids overly elongated shapes it allows for stronger deformation within its distance bounds.

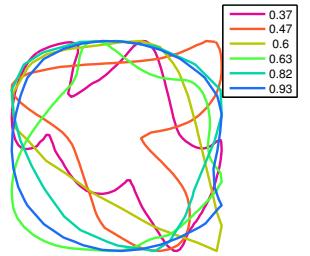


Figure 1.7: Different values of the isoperimetric quotient.

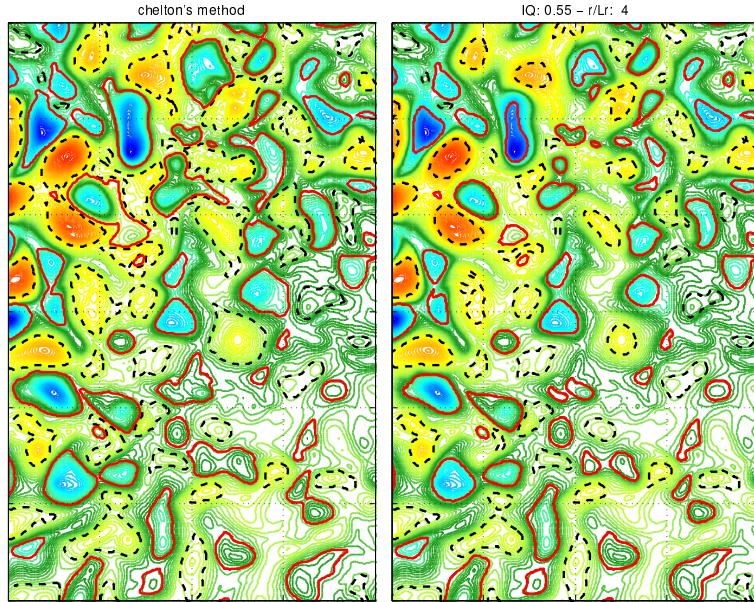


Figure 1.8: Left: The Chelton-method expects to detect eddies at their base and is rather tolerant with respect to the shape of found contour. The IQ-method aims more at detecting single round vortices without expecting found contour to be necessarily related to any howsoever-defined outer edge of the eddy.

Contour filter 9 Amplitude

```
function CR_AmpPeak
```

This function determines the amplitude *i.e.* the maximum of the absolute difference between **SSH** and current contour level and the position thereof as well as the amplitude relative to the mean **SSH** value of the eddy interior as done by Chelton *et al.* (2011). The amplitude is then tested against the user-given threshold. The function also creates a matrix with current contour level shifted to zero and all values outside of the eddy set to zero as well.

Contour filter 10 Chelton's Scales

```
function cheltStuff
```

Chelton *et al.* (2011) introduced 4 different eddy-scales.

1. The effective scale L_{eff} as the radius of a circle with its area equal to that enclosed by the contour.

2. The scale L_e as the radius at $z = e^{-1}a$ with a as the amplitude with reference to the original contour and the z -axis zero-shifted to that contour. In other words the effective scale of the contour that is calculated at $1/e$ of the original amplitude.
3. The scale $L = L_e/\sqrt{2}$.
4. The scale L_s which is a direct estimate based on the contour of SSH within the eddy interior around which the average geostrophic speed is maximum (Chelton et al., 2011). It is hence conceptually the same as σ . This scale was not calculated here, as I could not think of an efficient, simple way to estimate the area bounded by maximum geostrophic speed i.e. the zero-vorticity contour. To understand why this would be difficult to achieve see also contour-filters 11 and 12 and ??.

Contour filter 11 Profiles

`function EddyProfiles`

This step

- saves the meridional and zonal profiles of SSH, U and V through the Eddy's peak, spanning the entire sub-domain as described in contour-filter 4.
- creates spline functions from the ssh-profiles and uses them to interpolate the profiles onto 100-piece equi-distant coordinate vectors to build smooth interpolated versions of ssh-profiles in both directions.
- in turn uses the splined data to create smooth 4-term Fourier series functions for the profiles.

Contour filter 12 Dynamic Scale (σ)

`function EddyRadiusFromUV`

The contour line that is being used to detect the eddy is not necessarily a good measure of the eddy's scale i.e. it doesn't necessarily represent the eddy's outline very well. This becomes obvious when the area, as inferred by contour-filter 6, is plotted over time for an already successfully tracked eddy. The result is not a smooth curve at all. This is so because at different time steps the eddy usually gets detected at different contour levels. Since its surrounding changes continuously and since the eddy complies with the testing-criteria the better the closer the algorithm gets to the eddy's peak value, the determined area of the contour jumps considerably between time steps. This is especially so for large flat eddies with amplitudes on the order of 1cm. If the contour increment is on that scale as well, the difference in contour-area between two time steps

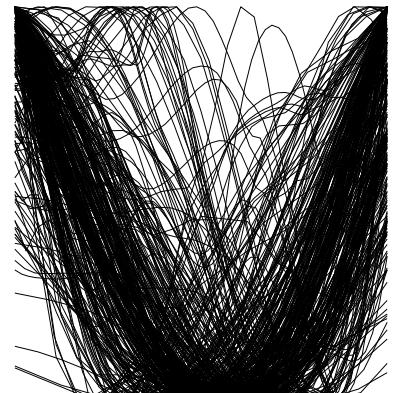
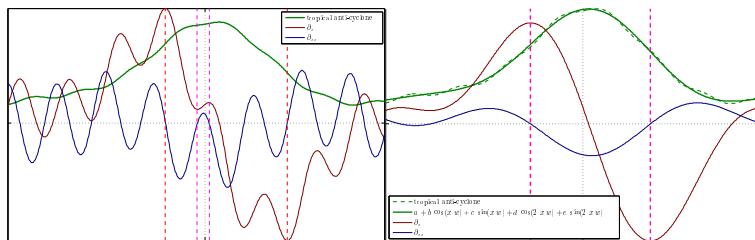


Figure 1.9: Zonal x - and z -normalized cyclone-profiles (early data $\sim 13/12$).

easily surpasses 100% and more. Since there is no definition for the *edge* of an eddy, it is defined here as the ellipse resulting from the meridional [zonal] diameters that are the distances between the first local extrema of orbital velocity (one negative, one positive) away from the eddy's peak in y - [x -] direction⁴. In the case of a meandering jet with a maximum flow speed at its center, that is shedding off an eddy, this scale corresponds to half the distance between two opposing center-points of the meander. It is also the distance at which a change in vorticity-polarity occurs and is thus assumed to be the most plausible dividing line between vortices. Trying to determine the location where this sign change in vorticity occurs in the profiles turns out to be very tricky. What we seek are local extrema of the geostrophic speeds *i.e.* of the ssh-gradients h_x . In a perfect Gaussian-shaped eddy, these would simply correspond to the first local extrema of h_x away from the peak. In *reality* the eddies can be very wobbly with numerous local maxima and minima in the gradients of their flanks. One could argue, that it must be the largest extrema, as it is the highest geostrophic speeds that are sought. In practice⁵ multiple superimposed signals of different scales often create very strong gradients locally. But the main issue here is that one weak eddy adjacent to one strong eddy also has the stronger gradients of the stronger one within its domain so that simply looking for the fastest flow speeds along the profiles is insufficient. It is also not possible to restrict the cut domain to the extent of a single eddy only, because at the time when the domain is selected, we do not know yet whether the detection algorithm *took bait* at the eddy's base or later close to the tip.

The best method thus far seems to be to use the Fourier-series functions from contour-filter 11 to determine the first extrema away from the eddy's peak (see fig. 1.10). The Fourier order was chosen to be 4 by trial and error. The effect is that small-scale low-amplitude noise is avoided, allowing for more reliable determinations of $\nabla^2 h_{\text{fourier}} = 0$.

Once the zero crossings in all 4 directions are found, their mean is taken as the eddy's scale (σ).



⁴ The velocities are calculated from the gradients of 4th-order Fourier fits to the SSH profile in respective direction (see contour-filter 11).

⁵ Especially for the high-resolution model data.

Figure 1.10: A flat wobbly low-latitude eddy resulting in multiple zero-crossings of its ∇^2 . The problem is addressed by differentiating the profile's Fourier-Series fit instead.

Contour filter 13 Dynamic Amplitude

```
function EddyAmp2Ellipse
```

As mentioned above, the contour that helps to detect the eddy is not representative of its extent. This is also true for the z -direction, for the same reasons. This function therefor takes an SSH-mean at indices of the ellipse created by the determined zonal and meridional *dynamical* diameters, and uses this as the basal value to determine a *dynamic* amplitude.

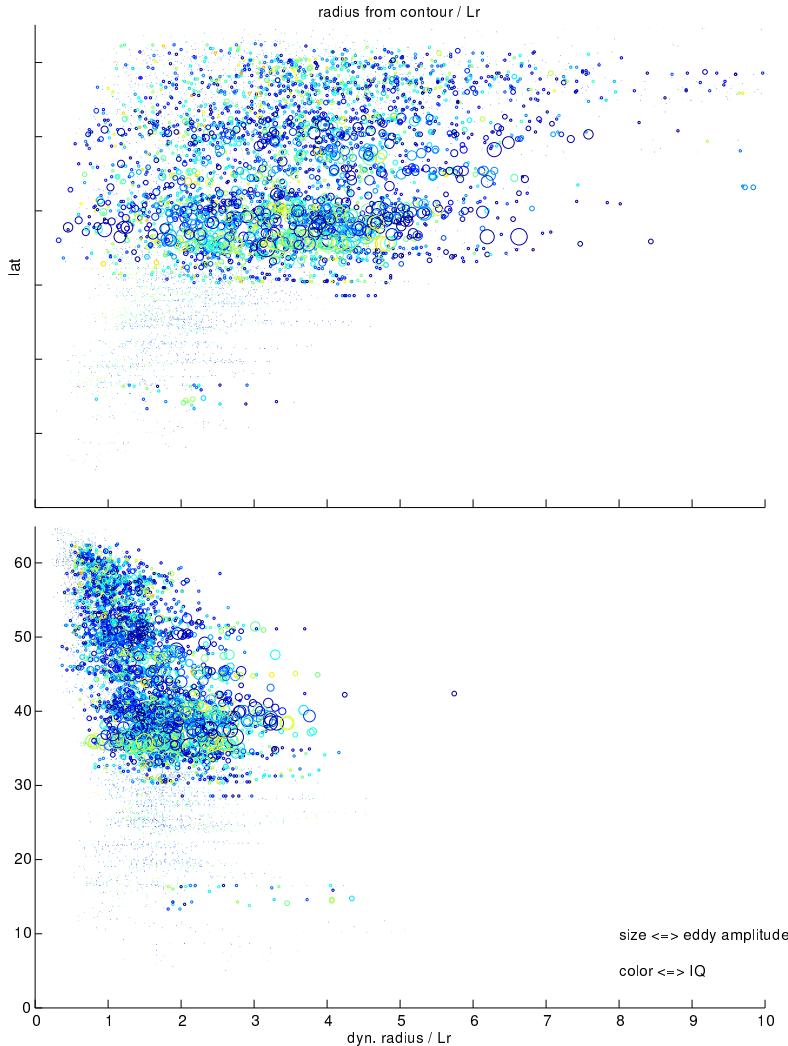


Figure 1.11: Eddies in the North-Atlantic. Y-axis: latitude. X-axis top: ratio of radius of circle with equal area to that of found contour to local Rossby-radius. X-axis bottom: ratio of σ to local Rossby-radius. Color-axis: Isoperimetric Quotient. Size: amplitude. The bottom plot suggests that a ratio of say 4 for σ/L_R^1 should be a reasonable threshold. Same graph for the Southern Ocean looks very different though (not shown here), in that said ratio often exceeds ratios as high as 10 and larger in the far south where L_R^1 becomes very small. This problem was addressed by prescribing a minimum value $L_R^1 = 20\text{km}$ for the calculation of the scale-threshold.

Contour filter 14 Center of Volume (CoV)

```
function CenterOfVolume
```

Instead of using the geo-position of the eddy's peak in the tracking procedure, it was decided to instead use the center of the volume created by the basal shifted matrix from contour-filter 9 i.e. the center of volume of the dome (resp. valley) created by capping off the eddy at the contour level. This method was chosen because from looking at animations of the tracking procedure it became apparent that, still using peaks as reference points,

the eddy sometimes jumped considerably from one time step to the next if two local maxima existed within the eddy. E.g. in one time-step local maximum A might be just a little bit larger than local maximum B and one time-step later a slight shift of mass pushes local maximum B in pole position, creating a substantial jump in the eddy-identifying geo-position hence complicating the tracking procedure.

Contour filter 15 Geo Projection

```
function ProjectedLocations
```

An optional threshold on the distance an eddy is allowed to travel over one time-step is implemented in the tracking algorithm in section 1.4. This is a direct adaptation of the ellipse-based constraint described by [Chelton et al. \(2011\)](#). The maximum distance in western direction traveled by the eddy within one time-step is limited according to $x_{west} = \alpha c \delta t$ with c as the local long-Rossby-wave phase-speed and

e.g. $\alpha = 1.75$. In eastern direction the maximum is fixed to a value of e.g. $x_{east} = 150\text{km}$. This value is also used to put a lower bound on x_{west} and for half the minor axis (y -direction) of the resultant ellipse.

This function builds a mask of eligible geo-coordinates with respect to the next time-step.

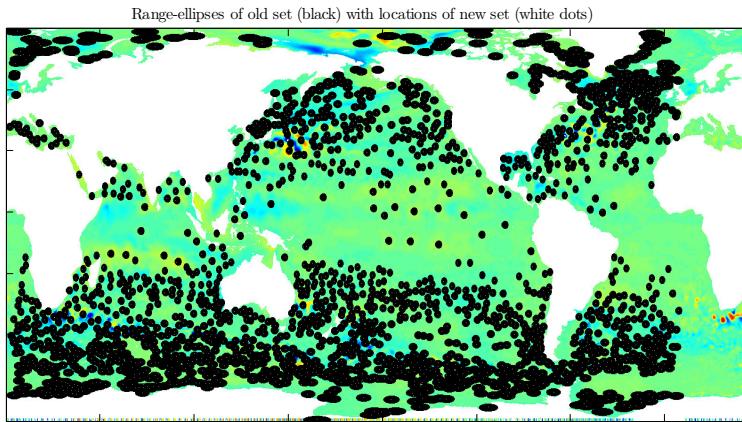


Figure 1.12: Among the saved meta-information for each eddy are also the indices describing the ellipse that defines the eddy's allowed locations for the next time-step.

1.5 Step S05: Track Eddies

```
S05_track_eddies
```

1.5.1 Main Tracking Procedure

DUE to the the relatively fine temporal resolution (daily) of the model data, the tracking procedure for this case turns out to be much simpler than the one described by [Chelton et al.](#).

(2007). There is almost no need to project the new position of an eddy, as it generally does not travel further than its own scale in one day. This means that one eddy can usually⁶ be tracked unambiguously from one time step to the next as long both time-steps agree on which eddy from the *other* time-step is located the least distance away. The algorithm therefore simply builds an arc-length-distance matrix between all old and all new eddies and then determines the minima of that matrix in both directions *i.e.* one array for the new with respect to the old, and one for the old with respect to the new set. This leads to the following possible situations:

- Old and new agree on a pair. *I.e.* old eddy O_a has a closest neighbor N_a in the new set and N_a agrees that O_a is the closest eddy from the old set. Hence the eddy is tracked. N_a is O_a at a later time.
- N_a claims O_a to be the closest, but N_b makes the same claim. *I.e.* two eddies from the new set claim one eddy from the old set to be the closest. In this situation the closer one is decided to be the old one at a later time-step and the other one must be a newly formed eddy.
- At this point all new eddies are either allocated to their respective old eddies or assumed to be *newly born*. The only eddies that have not been taken care of are those from the old set, that *lost* ambiguity claims to another old eddy, that was closer to the same claimed new eddy. *I.e.* there is no respective new eddy available which must mean that the eddy just *died*. In this case the entire track with all the information for each time step is archived as long as the track-length meets the threshold criterion. If it doesn't, the track is abandoned.

1.5.2 Improvements

The former is the core of the tracking algorithm. It is almost sufficient by itself as long as the temporal resolution is fine enough. The larger the time-step, the more ambiguities arise, which are attempted to be mitigated by flagging elements of the distance matrix not meeting certain thresholds:

- `function checkDynamicIdentity`
Consider the ambiguous case when there are two new eddies N_a and N_b in sufficient proximity to old eddy O_a . Let's assume O_a is a relatively solid eddy of rel. large scale with a steep slope *i.e.* large amplitude and that N_a is merely a subtle blob of an eddy whilst N_b is somewhat similar to O_a but with only half the amplitude. The situation then is clear: N_b is the, apparently slowly dying, O_a at a later time,

⁶ The only exception being the situation when one eddy fades and another emerges simultaneously and in sufficient proximity.

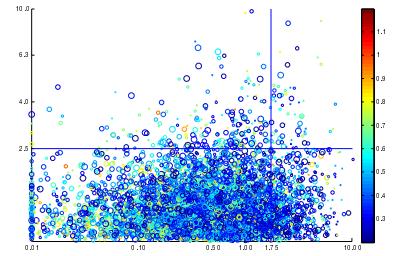
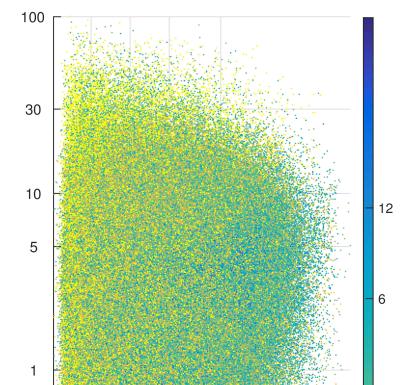


Figure 1.13: Each circle represents one eddy in the new time step. Y-axis: Maximum ratio to closest eddy in old set of either amplitude or σ , where 1 means *identical* and 2 means factor 2 difference. The threshold used for the final runs was 2. X-axis: Ratio of distance to closest eddy from old set divided by δt to local long-Rossby-wave phase-speed. Color-axis: Isoperimetric Quotient. Radius of circles: ratio of σ to local Rossby-radius. All eddies with said ratio larger than 10 are omitted. Note the obvious inverse correlation of scale to IQ, suggesting that all large eddies likely represent more than one vortex.



while N_a could either be a newly formed eddy, an old eddy with its respective representation in the old set something other than O_a , or even just temporary coincidental noise not representative of any significant mesoscale vortex at all. This interpretation should hold even when O_a sat right between the other two, thereby being much closer to O_a than N_b was. The purpose of this step is to make such decisions. It does so by comparing the *dynamic* versions of amplitude and scale (*ampToEllipse* and σ) between the time-steps. If either ratio from new to old⁷ surpasses a given threshold, the pair is flagged as non-eligible. It is important to use the *dynamic* parameters rather than those stemming from the contour line, because as mentioned in contour-filter 12, the contour line itself and the eddy's geometric *character* are hardly correlated at all. One eddy can get detected at different z -levels from one time-step to the next, resulting in completely different amplitudes, scales and shapes with respect to the contour.

The initial idea was, by assuming Gaussian shapes, to construct a single dimensionless number representing an eddy's geometrical character built upon the contour-related amplitude- and scale values only. Since we have no information about the vertical position of a given contour with respect to assumed Gauss bell, this problem turned out to be intrinsically under-determined and hence useless. The method eventually used, which checks amplitude and scale separately is again very similar to that described by Chelton *et al.* (see Box ??).

- `function` `nanOutOfBounds`

This is the second half of the prognostic procedure described in section 1.5. It simply flags all pairs of the distance matrix for which the index representing the *new* eddy's geographic location is not among the set of indices describing the ellipse⁸ around respective *old* eddy.

⁷ In order to compare in both directions equally: $\exp(|\log(v_n/v_o)|)$ where v is either amplitude or scale.

- `function` `checkAmpAreaBounds`

This is the direct adaptation of Chelton *et al.*'s description of how to test for sufficient similarity of amplitude and area between time steps.

⁸ see figure fig. 1.12.

Step S06: Cross Reference Old to New Indices

```
function S05_init_output_maps
```

The output Mercator-maps usually have different geometry from the input maps'. This step allocates all grid nodes of the input data to their respective nodes in the output map. Each output cell will then represent a mean of all input-nodes falling into that quadrilateral.

1.6 *Running the Code*

THE separate steps can be run all at once (`sall.m`) or one by one, as long as they are started consecutively in the order indicated by their name (`s00..`, then `s01..` etc.). `s01b` is not necessary though. Each step either creates its own output files or extends old ones, which are then read by the next step. All output data is saved in the user given root-path. This concept uses quite a lot of disk space and is also slowed substantially by all the reading and writing procedures. The benefit is that debugging becomes much easier. If the code fails at some step, at least all the calculations up to that step are dealt with and do not need to be re-run. The concept also makes it easier to extend the code by further add-ons.

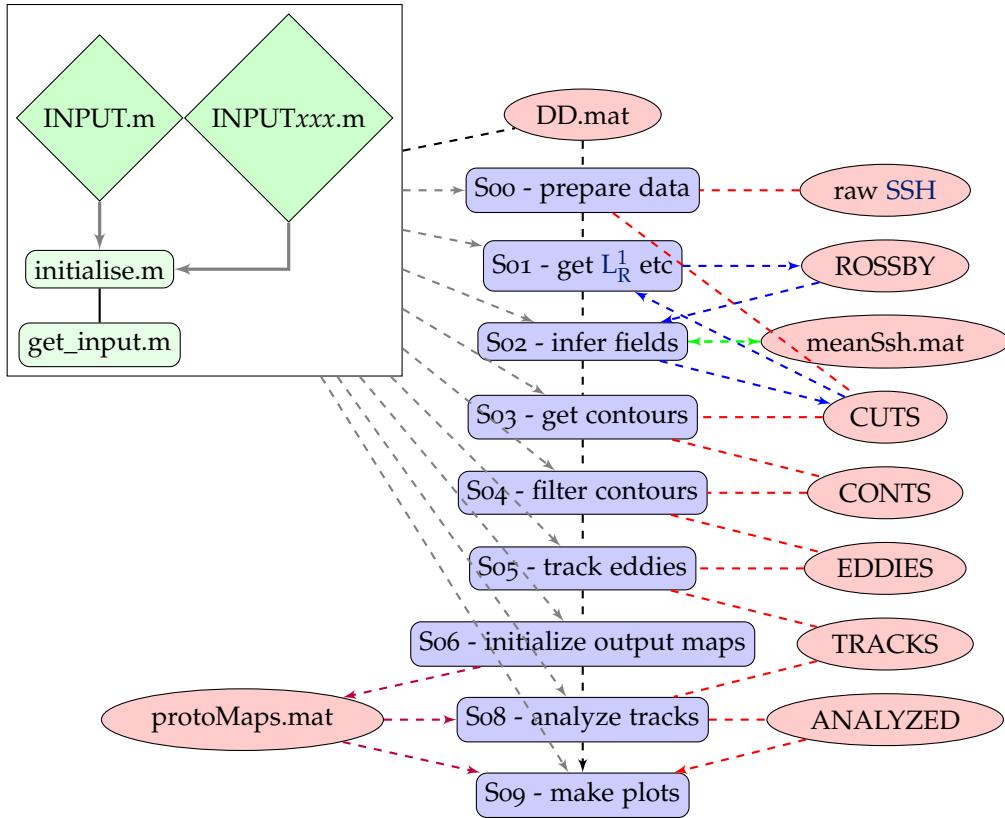


Figure 1.15: Basic code structure. The only files that are to be edited are the INPUT files. INPUT.m is independent of the origin of data, whilst the files INPUTaviso.m, INPUTpop.m etc set are samples of more source-specific parameter settings. Each of the SXx-steps initially calls initialise.m, which in turn scans all available data, reads in the INPUT data via get_input.m, corrects for missing data etc and creates DD.mat. The latter is the main meta-data file, which gets updated throughout all steps. All data is built step-by-step along the consecutive SXx-steps (red line). The SXx-steps are the only programs that have to be called (in order) by the user. Beware that missing data (in time) is interpolated automatically in each step. Note also that meanSsh.mat should be recalculated if the time span is changed!